

# TDDI41 - Systeminstallation

Anders Fröberg

anders.froberg@liu.se

# Upplägg

- Vem är jag ?
- Kursen
  - Innehåll
  - Upplägg
- Intrduktion till Linux



# Kursinformation

- Kurshemsidan: <http://www.ida.liu.se/~TDDI41/>
  - Deadlines, labbeskrivningar, labanmälan, errata
- Laborationer
  - Studera – **planera** – implementera – **testa** – dokumentera
  - Krav på ”ingenjörsmässigt” arbetssätt
  - Strikta deadlines

# Kursens kunskapsfilosofi

**People generally remember...  
(learning activities)**

**10% of what they read**

**20% of what they hear**

**30% of what they see**

**50% of what they see and hear**

**70% of what they say and write**

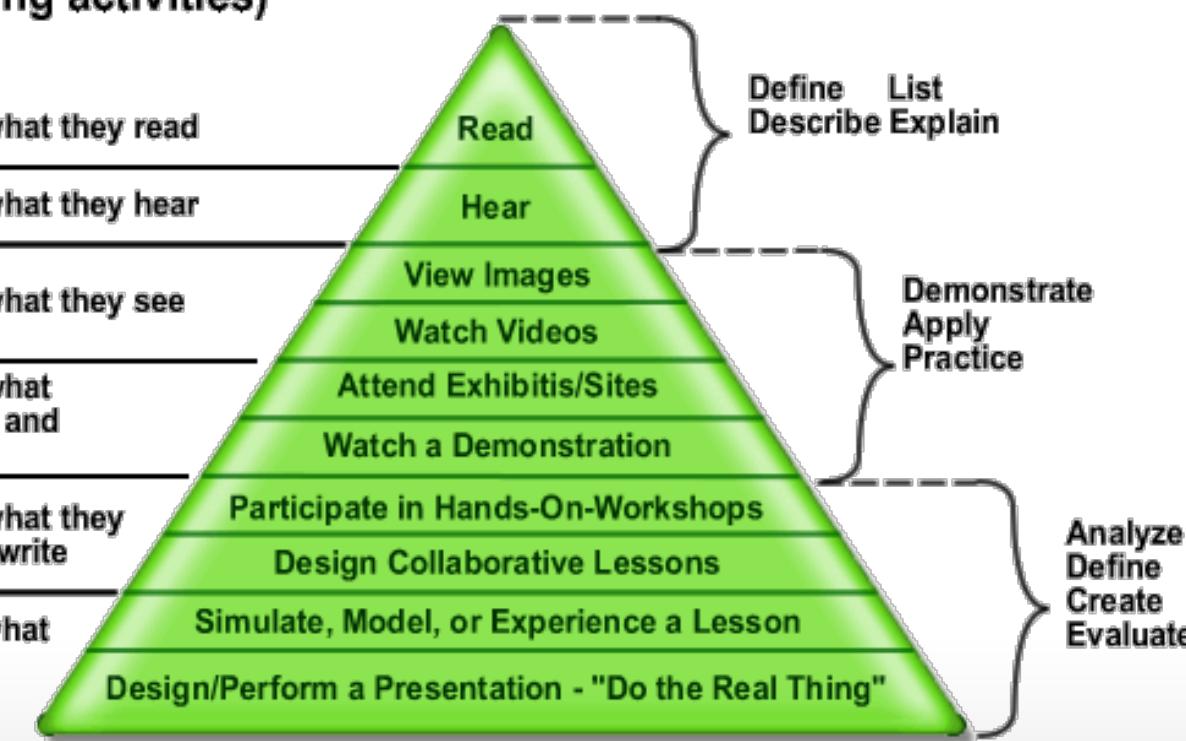
**90% of what they do.**

**People are able to...  
(learning outcomes)**

**Define List  
Describe Explain**

**Demonstrate  
Apply  
Practice**

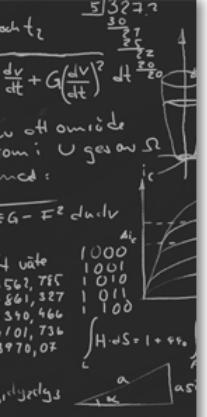
**Analyze  
Define  
Create  
Evaluate**



# Labmiljö

- Eget (virtuellt) labnät
- Alla på en server
- Virtualiserade system
  - User-mode Linux
  - Debian/Gnu Linux





<http://www.ida.liu.se/~TDDI41>

---



**MATURITY**  
has nothing to do with  
age. Maturity comes  
from **experiences**,  
**mistakes**, **learning**,  
and **understanding**.

[WWW.LIVELIFEHAPPY.COM](http://WWW.LIVELIFEHAPPY.COM)

Grade

A-

B

A



TEAM

# Kursens kunskapsfilosofi

**People generally remember...  
(learning activities)**

**10% of what they read**

**20% of what they hear**

**30% of what they see**

**50% of what they see and hear**

**70% of what they say and write**

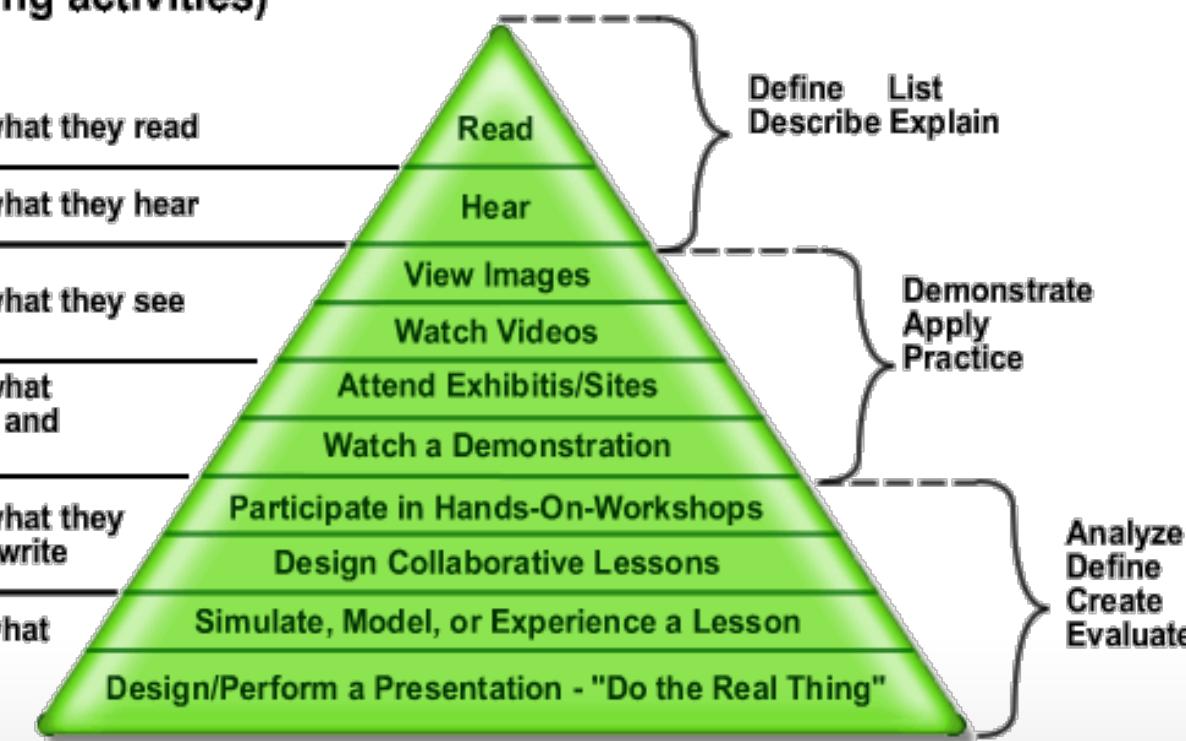
**90% of what they do.**

**People are able to...  
(learning outcomes)**

**Define List  
Describe Explain**

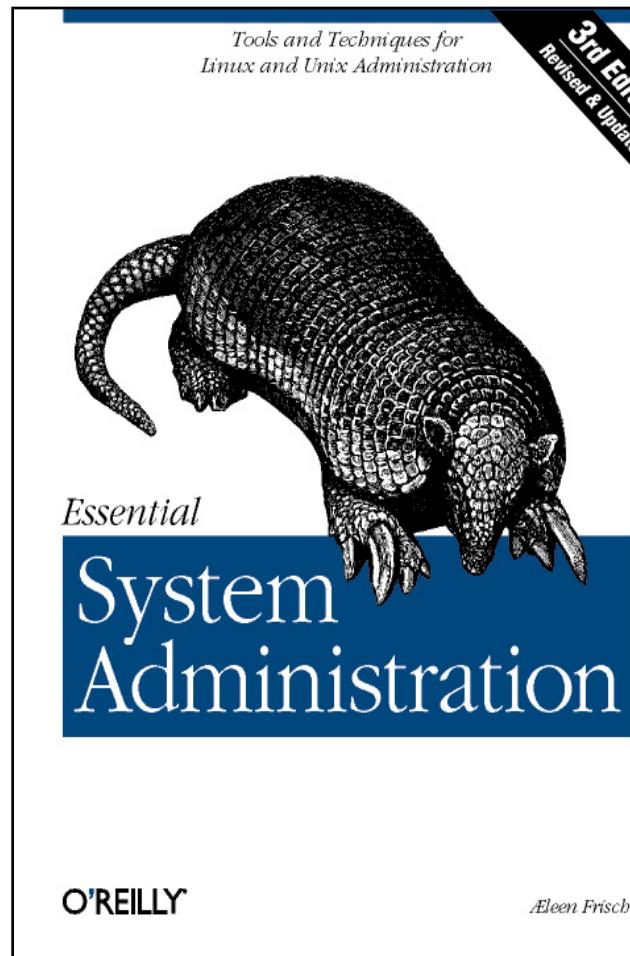
**Demonstrate  
Apply  
Practice**

**Analyze  
Define  
Create  
Evaluate**





# Rekommenderade böcker



**Essential System Administration, 3rd Ed.**  
*Aileen Frisch*

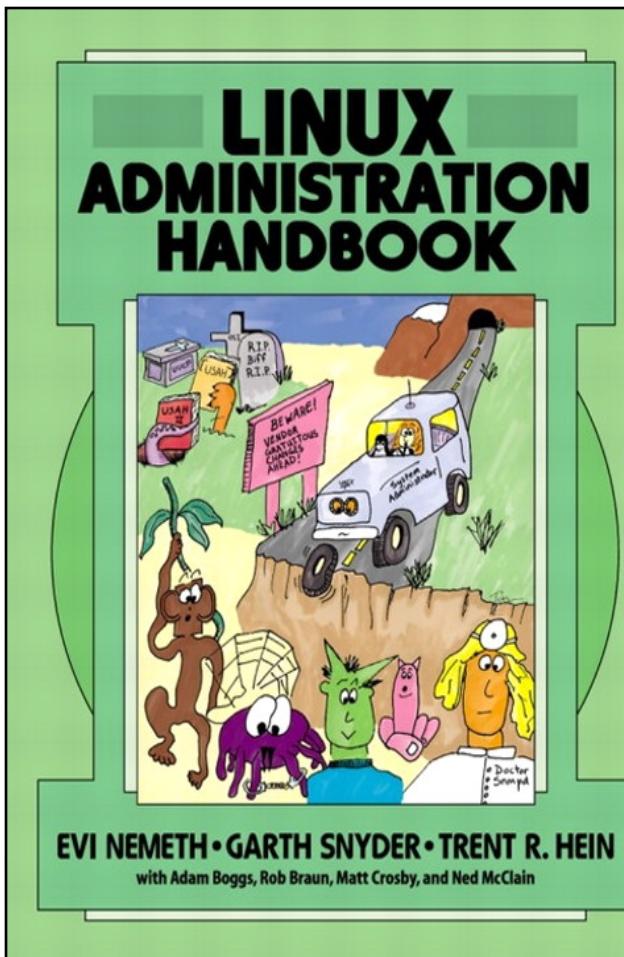
Grundlig genomgång av administration av Unix-system. Tar upp alla viktiga delsystem på ett bra sätt. Exempel från Linux, Tru64, BSD, HP-UX, AIX, Solaris och andra system. Fullständigt fokus på tekniken. En hård bok.

Aileen Frisch





# Rekommenderade böcker



## Linux Administration Handbook

*Evi Nemeth, et al*

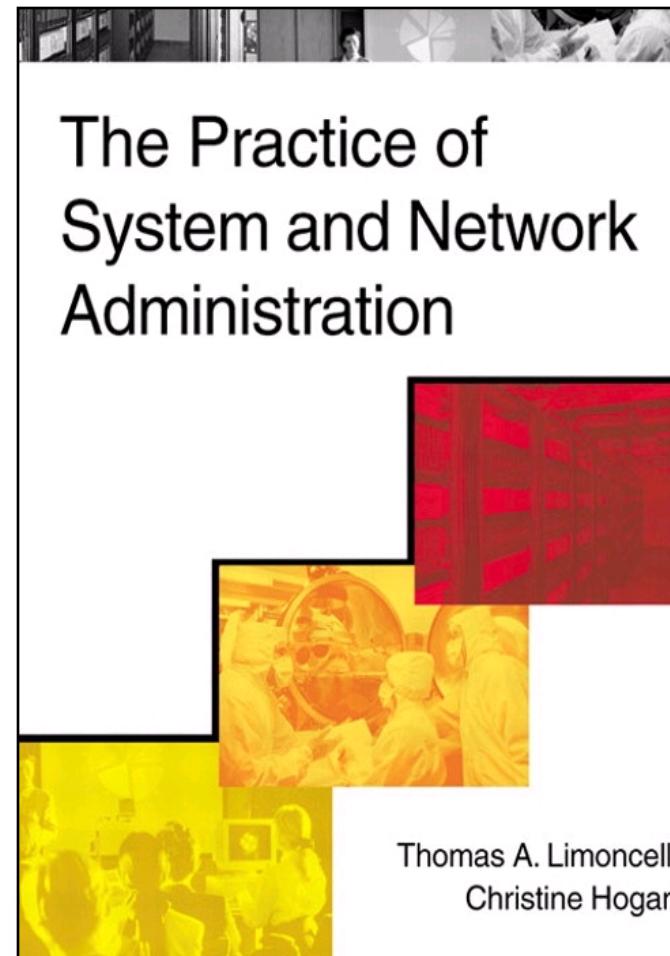
Administration av Linux-system. Täcker alla viktiga delsystem både ur Linux-perspektiv och mer generellt perspektiv. Likvärdig Essential System Administration men med ett något annorlunda urval av ämnen. En ganska hård bok.

Evi Nemeth





# Rekommenderade böcker



## The Practice of System and Network Administration

*Thomas A. Limoncelli*

*Christine Hogan*

Fokus på systemadministratörens yrke, inte på tekniken. Tar upp alla aspekter, från kundkontakter och grupp dynamik till hur man gör bra kabeldragning. Massivt innehåll. Mycket lite teknik. En mjuk bok.

Tom Limoncelli

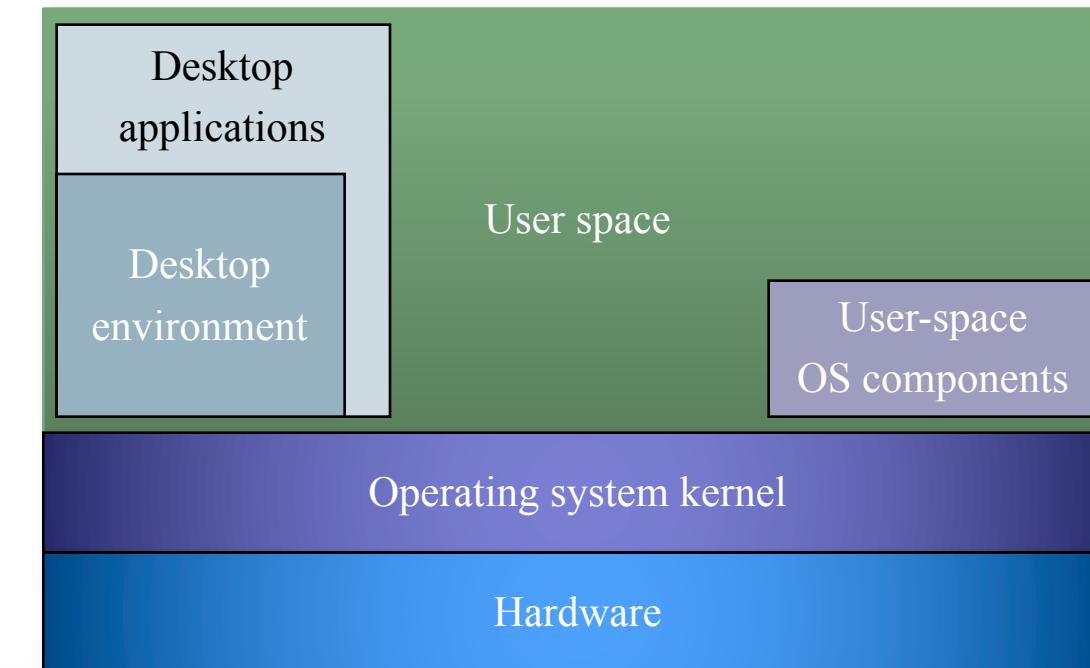


Christine Hogan

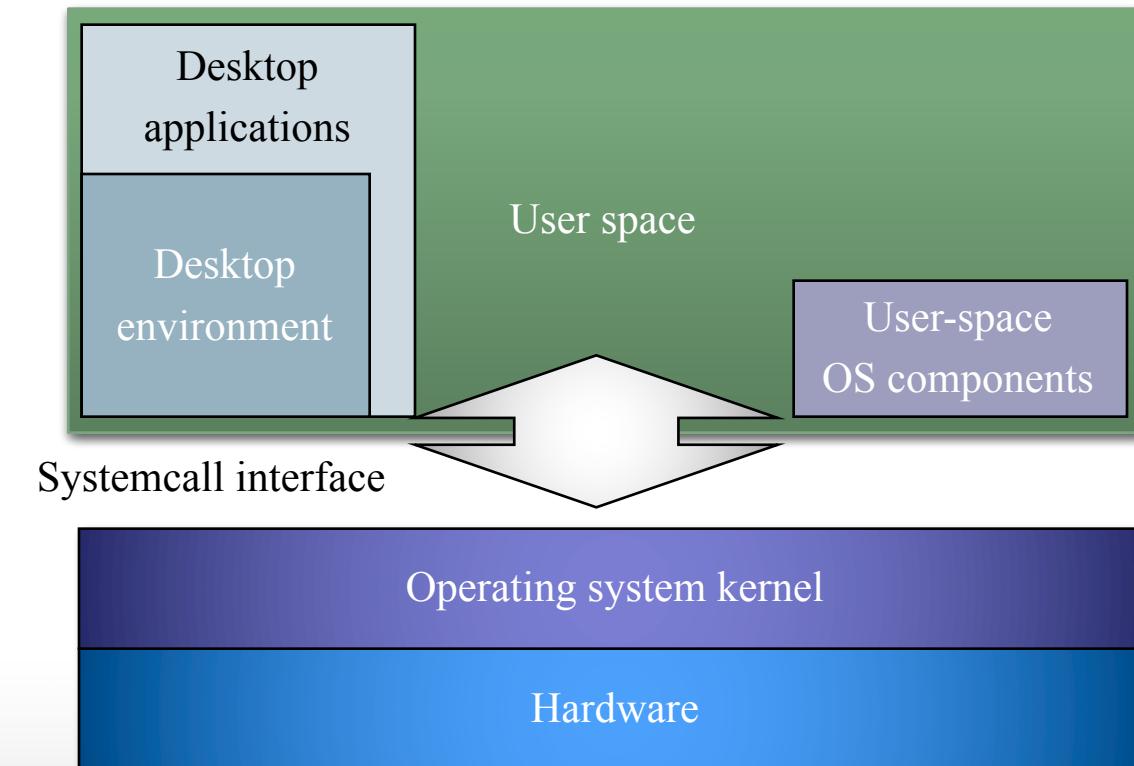


# LINUX

# What is Linux?



# What is Linux?



**1969**

**1970**

**1971**

**1972**

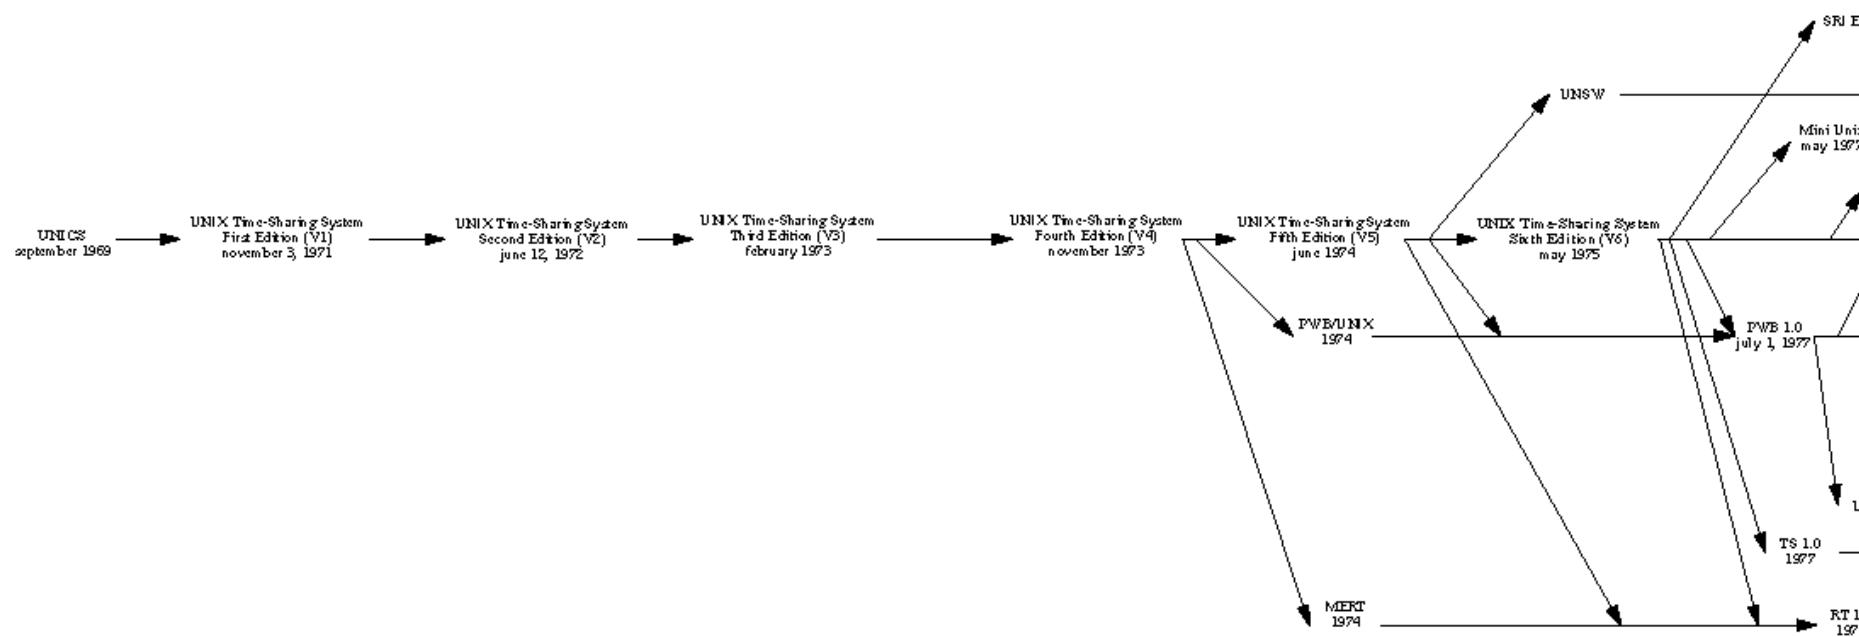
**1973**

**1974**

**1975**

**1976**

**1977**



# functions layers

## user space interfaces

## virtual subsystems

## transformations

## functional subsystems

## devices control

## hardware interfaces

## electronics

### system

#### system interfaces

```
linux/syscalls.h
asm-i386/uaccess.h
copy_from_user
cdev
cdev_map
sys_reboot
```

#### Device Model

```
drivers/base/
subsystem_register
subsystem
class
class_device
class_device_create
device_driver
probe
device
bus_type
device_create
driver_register
```

### networking

#### sockets access

```
sys_socketcall
sys_socket
socket_file_ops
```

#### protocol families

```
sock_create
socket
inet_family_ops
inet_create
unix_family_ops
proto_ops
inet_dgram_ops
inet_stream_ops
```

### storage

#### files & directories access

```
sys_open
sys_read
sys_write
do_path_lookup
sys_mount
```

#### Virtual File System

```
vfs_read
vfs_write
inode
file
vfs_create
file_operations
file_system_type
get_sb \ super_block
```

### memory

#### memory access

```
sys_brk
sys_mmap2 /proc/self/maps
```

#### Virtual memory

```
virtually contiguous memory
vmalloc
vmlist
vm_struct
```

### processing

#### Processes

```
sys_execve
fs/exec.c
linux_binfmt
sys_nanosleep
```

#### threads

```
kernel_thread
semaphore
msleep
wait_queue_head_t
work_struct
workqueue_struct
```

### human interface

#### HI char devices

```
cdev
/dev/input/mice
stdin
stdout
/dev/snd/...
/dev/dsp
/kmsg
```

```
printk
```

## transformations

## functional subsystems

## generic HW access

### devices access and bus drivers

```
usb_driver
pci_driver
usb_hcd
request_region
request_mem_region
ioremap
```

#### protocols

```
proto
udp_prot
tcp_prot
ip_rcv
ip_queue_xmit
ip_forward
```

#### virtual network device

```
net_device
neif_rx
dev_queue_xmit
alloc_etherdev
ether_setup
alloc_ieee80211
ieee80211_rx
ieee80211_xmit
```

#### Logical File Systems

```
ext3_file_operations
ext3_get_sb
```

#### Block devices

```
block/
gendisk
block_device_operations
init_scsi
scsi_device
scsi_driver
sd_fops
usb_storage_driver
```

### disk controllers drivers

```
Scsi_Host
aic94xx_init
ide_disk_ops
ide_intr
ide_do_request
```

#### logical memory

```
mm/slab.c
/proc/slabinfo
kmalloc
kmem_cache
slab
```

#### Page Allocator

```
mm/page_alloc.c
/proc/buddyinfo
get_free_pages
_alloc_pages
pgdat_list
```

#### Scheduler

```
kernel/sched.c
schedule_timeout
task_struct
process_timeout
activate_task
current ~> thread_info
```

#### interrupt context

```
timer_list
jiffies ++
setup_irq
timer_interrupt
do_IRQ
do_softirq
tasklet_struct
```

#### HI subsystems

```
tty
drivers/media/
video dev init
alsa oss sound/
```

#### abstract devices and HID class drivers

```
console_fops
console
kbd
mousedev
fb_fops
```

## network devices drivers

### network devices drivers

```
e100_open
e100_open
rt8139_open
ipw2100_open
zd1201_net_open
```

### disk controllers

```
SCSI
IDE
SATA
```

### physical memory operations

```
MMU
RAM
pte
page
do_page_fault
```

### CPU specific

```
arch/i386/kernel/
/proc/interrupts
system_call
show_regs
trap_init
pt_regs
cli
sti
```

### HI peripherals device drivers

```
atkbd
psmouse
drivers/video/
i8042_driver
ac97_driver
drivers/media/
```

## electronics

### network controllers

```
Ethernet
WiFi
```

### disk controllers

```
SCSI
IDE
SATA
```

### memory

```
MMU
RAM
```

### CPU

```
registers
APIC
interrupt controller
```

### user peripherals

```
keyboard
graphics card
audio
```

# Linux distributions

- Distributions offer a complete operating environment
  - Installation software/procedures
  - Operating system kernel
  - Complete user space
  
- Different distributions offer different choices
  - Philosophy (e.g. open/closed; stable/cutting-edge; source/binary)
  - Choice of user space software (e.g. desktop environment)
  - Purpose (e.g. desktop/server)
  - Size (e.g. large/small disk; large/small memory)
  - Special-purpose (e.g. PVR, forensics, pen-test)

1991 1992 1993 1994 1995 1996 1997 1998 1999 2000 2001 2002 2003 2004 2005 2006 2007

# Linux distro timeline

Version 7.2 by NPU (nonplusx@gmail.com)

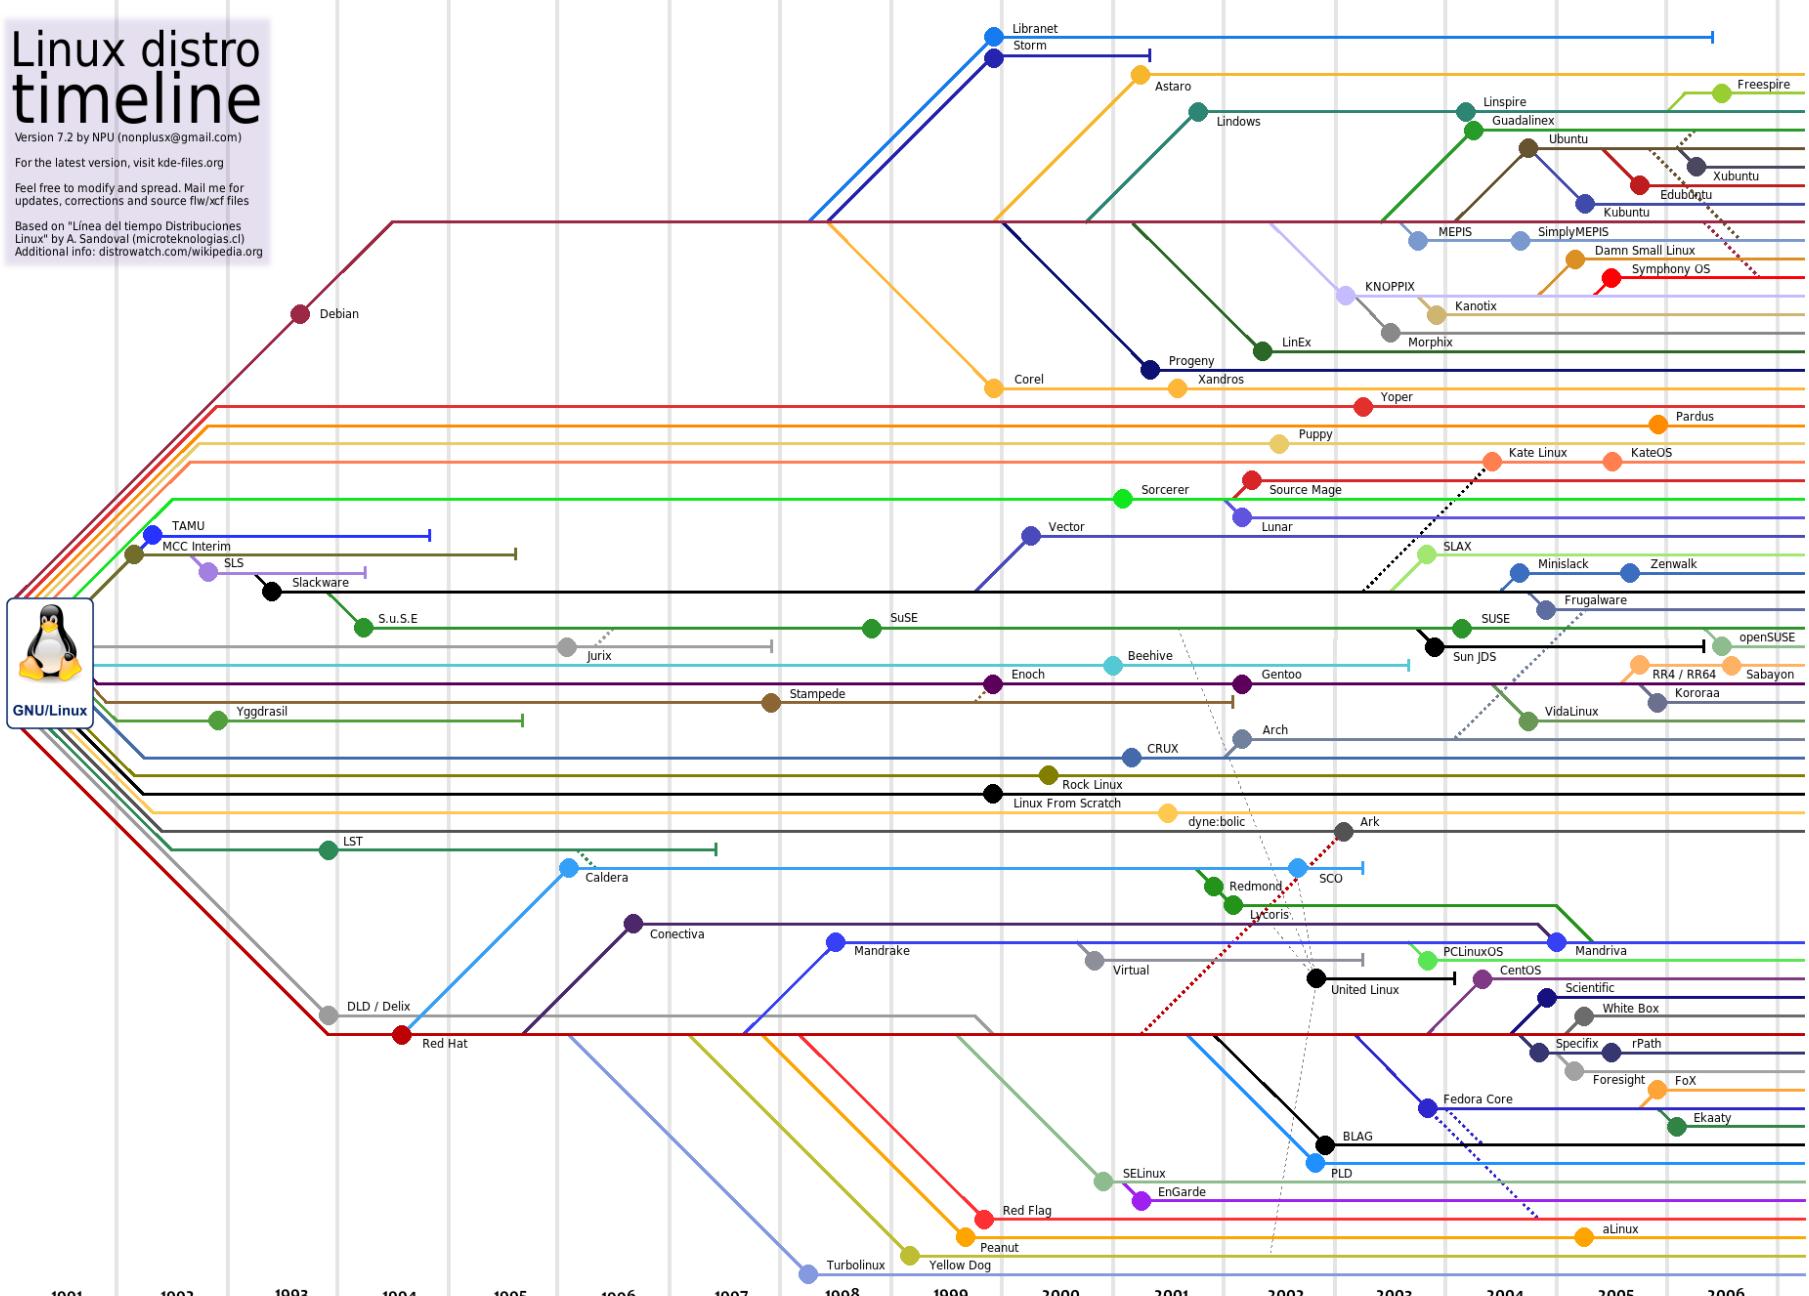
For the latest version, visit kde-files.org

Please feel free to modify and spread. Mail me for updates, corrections and source flw/xcf files

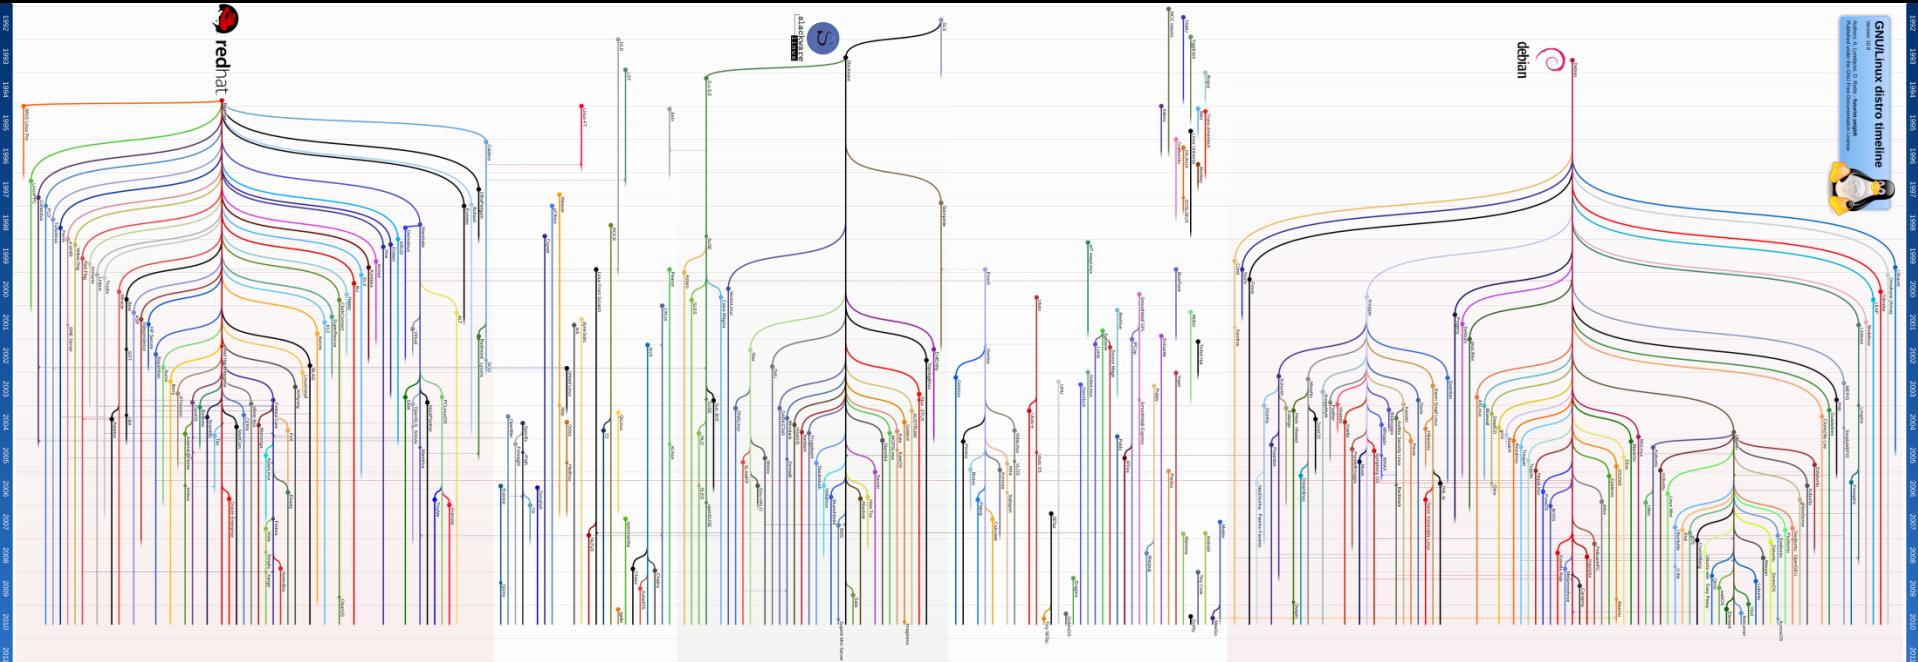
Based on "Línea del tiempo Distribuciones Linux" by A. Sandoval (microteknologias.cl)  
Additional info: distrowatch.com/wikipedia.org



GNU/Linux



# 1992



# 2010

GNU/Linux distro timeline  
Version 0.1.0  
Author: Antonio O. Salas  
Available under the CC-BY 3.0 Creative Commons License

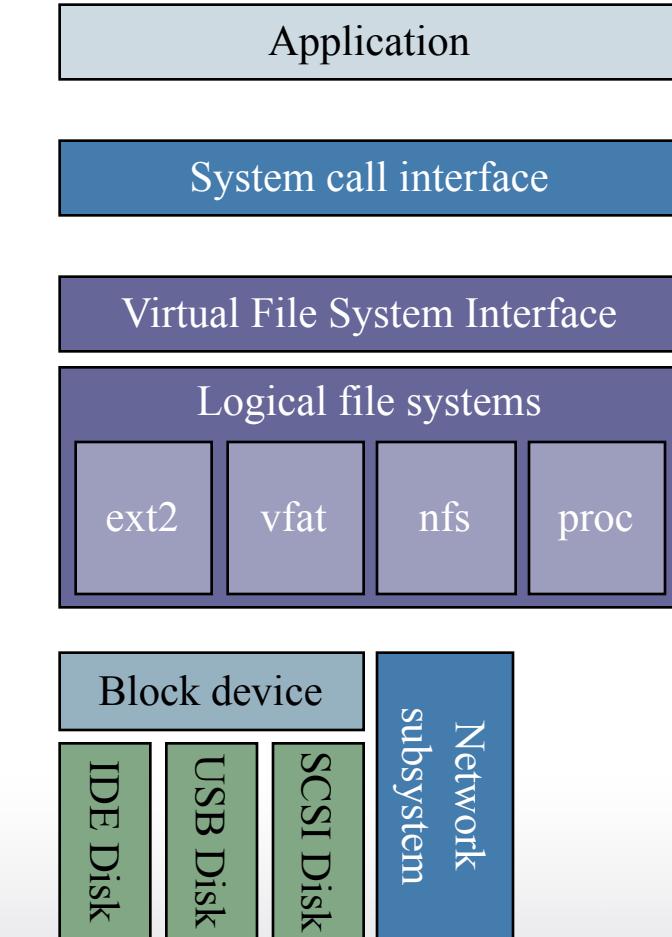


# What about Unix?

- What is Unix?
  - Unix was an operating system developed by Bell Labs, then AT&T, then many, many others
  - Unix is a trademark held by the Open Group
  - Unix is an *idea*
- Standard Unix: SUS and POSIX
- Linux *feels* like Unix
  - Same (or similar) user space
  - Based on similar ideas
    - Less now with systemd

# A look at Linux: files and file systems

- Applications go through VFS
- VFS passes all requests to appropriate logical file system
- Logical file system gets data from wherever it wants to
  
- Unified interface to all files
- *Anything* can be a file



# What do we mean when we say "file system"?

- We mean the software that VFS uses to get files and data
  - "Have you loaded ufs into the kernel?"
- We mean the specification of the data structures that the software uses to store files and data
  - "The ext2 superblock is replicated for reliability purposes."
- We mean the *actual* data structures stored on a disk (and the simulated ones provided by non-storage file systems)
  - "I need to verify the integrity of the file system on /dev/hda1."

# File system organization

- Single directory hierarchy
  - No "drive letters"
  - Starts at /
- "Mounting"
  - Attaches a file system (e.g. ext2 file system on a hard drive, or file system exported using NFS) to a point in the file system hierarchy
  - The "root file system" is the file system mounted at /, and must be mounted during system startup
  - It is possible to mount the same file system several times in different locations
  - It is possible to specify options (e.g. read only) when mounting
  - Default mounts in /etc/fstab (on Linux)

# Example of mounting file systems

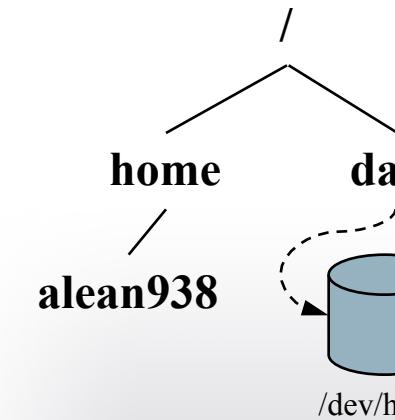
```
mount -t ext2 -o noexec /dev/hda2 /data  
mount -t nfs -o rw atlas:/users/9/3/alean938 /home/alean938
```

**File system type:** ext2

**Options:** noexec

**Source:** /dev/hda2 (blockenhet)

**Mount point:** /data

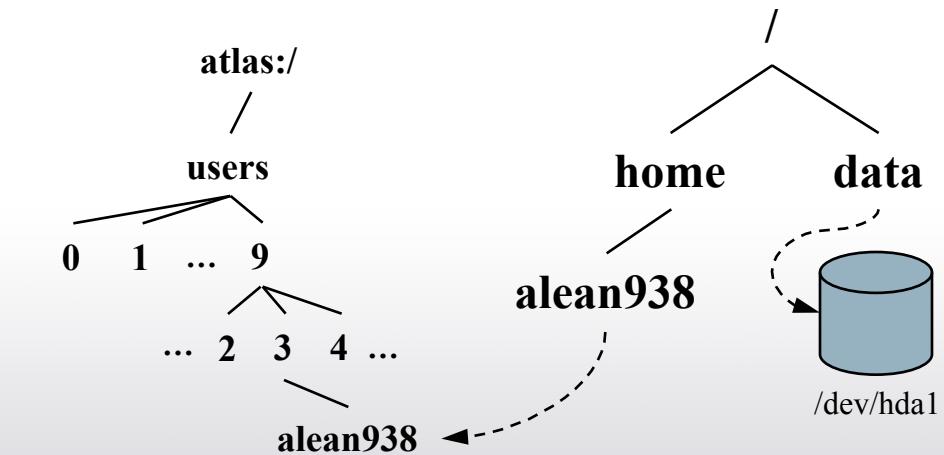


**File system type:** nfs

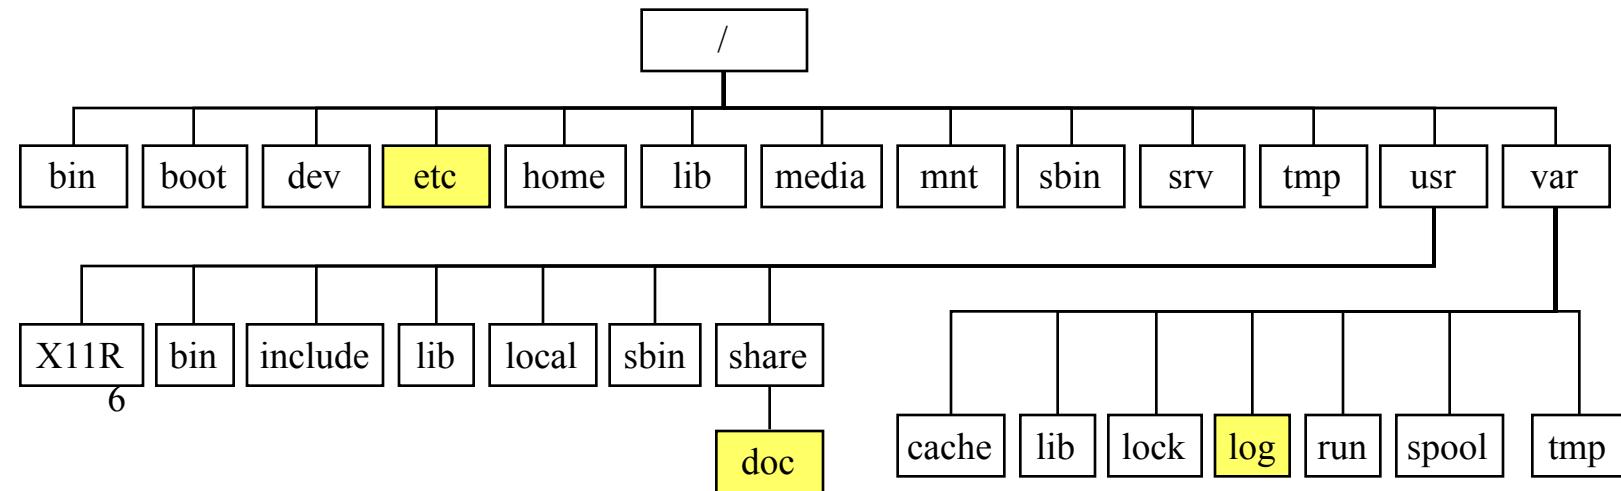
**Options:** rw

**Source:** atlas:/users/9/3/alean938

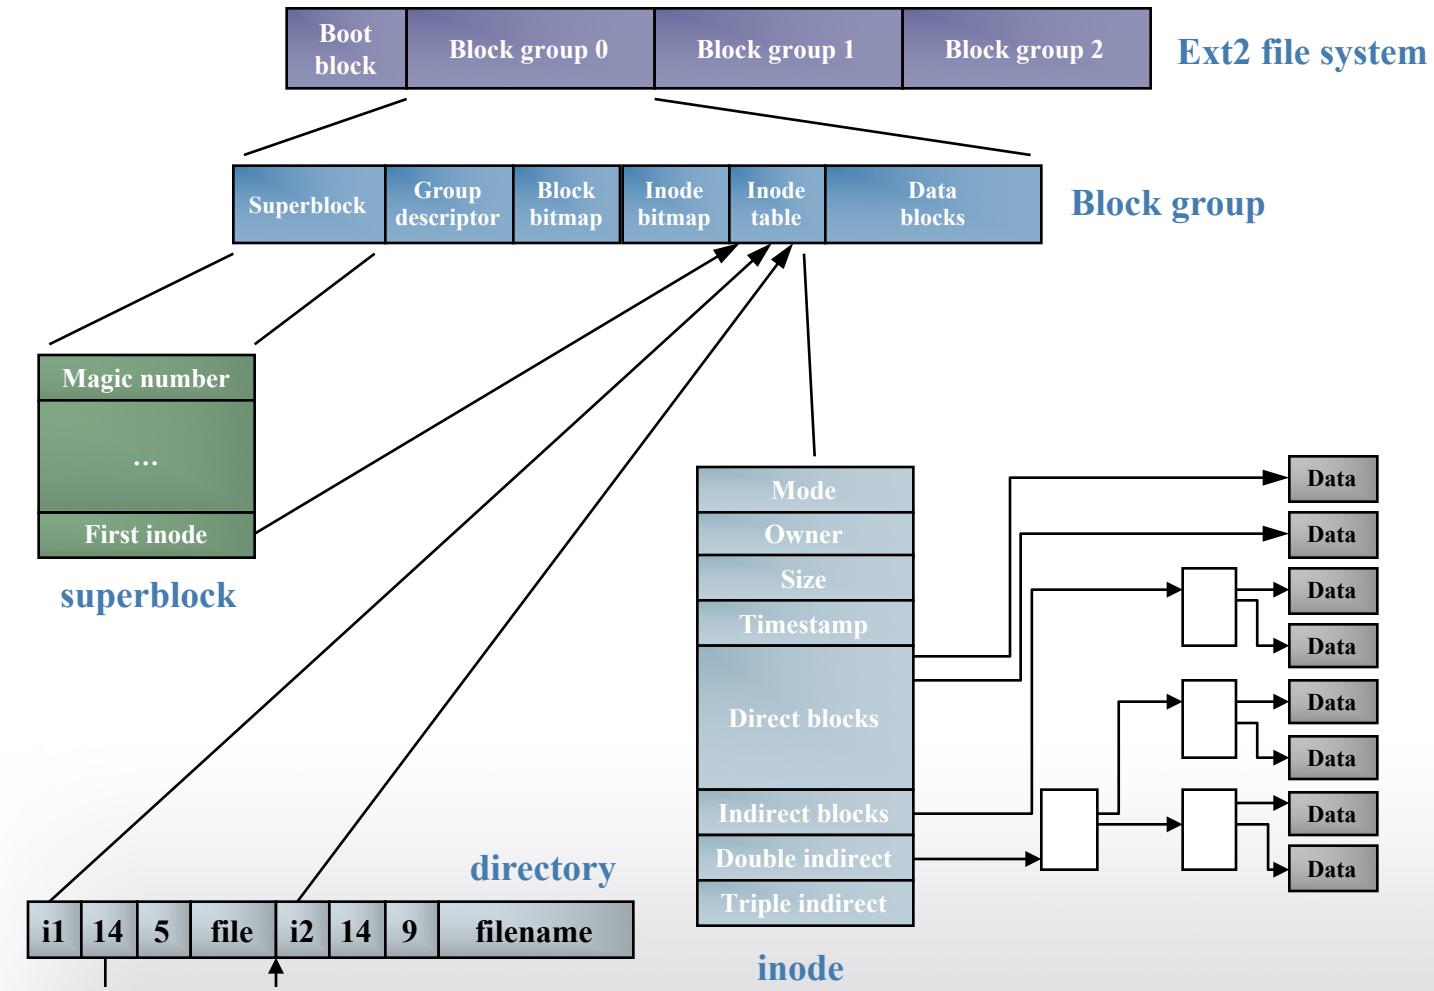
**Mount point:** /home/alean938



# Standard hierarchy



# Overview of the ext2 filesystem (on disk)



# The inode

- Every filesystem object is represented by an inode
  - The VFS does not care about different kinds of objects
  - Higher layers may treat different objects differently
- Some kinds of objects
  - Regular files
  - Directories
  - Device nodes
  - Sockets
  - Pipes
- Some contents of an inode
  - Creation and modification time
  - Mode (permissions)
  - Owner and group
  - Behavior (read, write, etc.)
- The inode does *not* contain the name or location of an object
- Names are *linked* to inodes through directory entries
  - Every inode can have multiple names (known as links)
  - An inode does not need any names at all

# An inode example: multiple links

- You can see inodes using the ls command in Linux
  - Use ls -i to see inode numbers
  - Use ls -l to see more data

Same inode but different names

```
% cd /bin  
% ls -i bz*  
349924 bzcat 349923 bzegrep 349923 bzgrep 349930 bzless  
349925 bzcmp 349928 bzexe 349924 bzip2 349930 bzmore  
349925 bzdiff 349923 bzfgrep 349931 bzip2recover  
% ls -li bzegrep bzfgrep bzgrep  
349923 -rwxr-xr-x 3 root root 3642 Aug 25 2006 bzegrep  
349923 -rwxr-xr-x 3 root root 3642 Aug 25 2006 bzfgrep  
349923 -rwxr-xr-x 3 root root 3642 Aug 25 2006 bzgrep
```

Mode  
Number of links  
Owner and group

# Another example of multiple links

- Directories have multiple links
  - One for the name
  - One for “.”
  - One for each “..” in each subdirectory

Just different names  
for the same inode

```
% ls -ldl /usr /usr/. /usr/bin/.. /usr/local/..
206753 /usr
206753 /usr/.
206753 /usr/bin/..
206753 /usr/local/..
% ls -ldi /usr
206753 drwxr-xr-x 13 root root 4096 Apr 20 12:06 /usr
%
```

One for /usr  
One for /usr/.  
Eleven for subdirectories

# Another inode example: special files

- Non-file and directory objects
  - Block devices ("block special"; e.g. hard disks)
  - Character devices ("character special"; e.g. sound cards)
  - Unix domain sockets
  - Pipes (FIFOs)

```
% ls -l /tmp/.X11-unix/X0
srwxrwxrwx 1 root root 0 Aug 13 13:28 /tmp/.X11-unix/X0
% ls -l /dev/hda1
brw-rw---- 1 root disk 3, 1 Aug 13 13:27 /dev/hda1
% mknod /tmp/fifo p
% ls -l /tmp/fifo
prw-r--r-- 1 root root 0 Aug 17 11:01 /tmp/fifo
% ls -l /dev/audio
crw-rw---- 1 root audio 14, 4 Aug 13 13:27 /dev/audio
%
```

Socket  
Block special  
Pipe (FIFO)  
Character special



# Permissions

% ls -l example

-rwxrwxrwx 1 root root 0 Aug 17 12:27 example

%

	Read	Write	Execute
User	4	2	1
Group	4	2	1
Other	4	2	1

Numerical permissions

User:      Read + Write + Execute                        =                     $4 + 2 + 1 = 7$

Group:     Read + Execute                                =                     $4 + 2 = 6$

Other:     Execute                                        =                     $1 = 1$

**Mode:**

**761**

# How do "user" and "group" work?

- Conceptually: if the user who owns the file attempts to access it, use "user" permissions; if a user who belongs to the same group as the file attempts to access it, use "group" permissions
- But how does the operating system know?
- All accesses are performed by processes and...
- ...every process has an "effective user ID" (EUID) and an "effective group ID" (EGID)
- ...every user belongs to a set of groups

# Where are users and groups defined

- Depends on what /etc/nsswitch.conf says
  - The "name service switch" allows users and groups (and other data) to be retrieved from many different sources
  - The default is "local files" (/etc/passwd and /etc/group)

```
% cat /etc/nsswitch.conf
# /etc/nsswitch.conf
passwd:      files nis
group:       files nis
shadow:      files
```

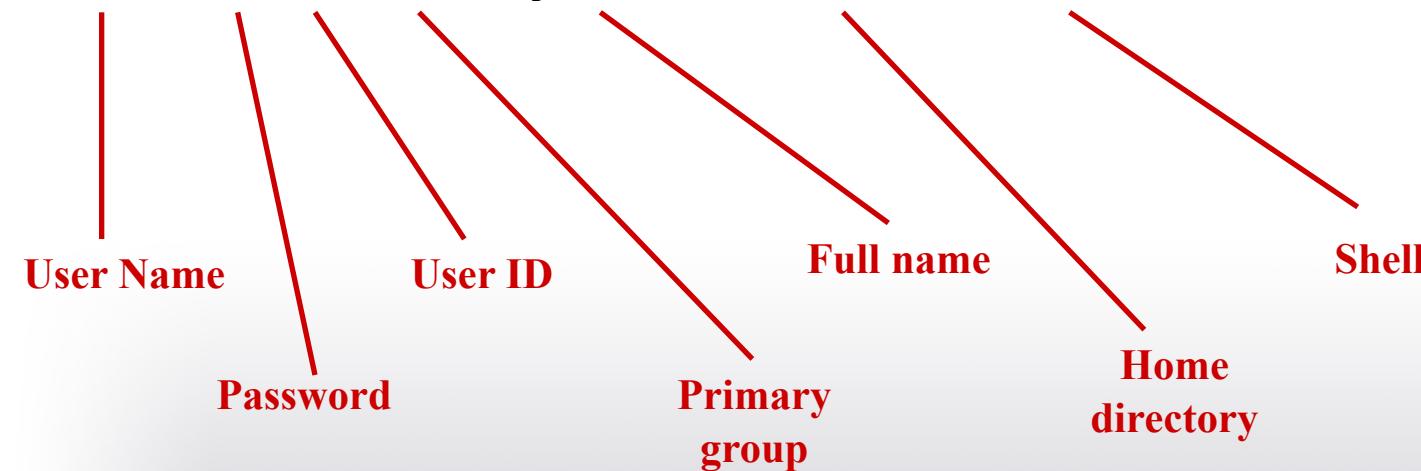
Users are stored in local files and in NIS

Groups are stored in local files and in NIS

# How are users defined?

```
% head -4 /etc/passwd
root:x:0:0:root:/root:/bin/bash
daemon:x:1:1:daemon:/usr/sbin:/bin/sh
bin:x:2:2:bin:/bin:/bin/sh
sys:x:3:3:sys:/dev:/bin/sh

% ypcat passwd | head -4
lindu932:*:24083:6000:Linus Dunkers:/home/lindu932:/bin/tcsh
frehi371:*:20883:6000:Frederik Hillerborg:/home/frehi371:/etc/notsignedsh
thogu383:*:17926:6000:Thomas Gustafsson:/home/thogu383:/etc/closedsh
tanmi350:*:55718:6000:Tanja Mitic:/home/tanmi350:/bin/tcsh
```



# How are groups defined?

```
% tail -4 /etc/group
```

```
ntp:x:111:
```

```
smmta:x:112:
```

```
smmsp:x:113:
```

```
uml-net:x:114:
```

```
% ypcat group | head -4
```

```
TDBD34:*:9197:eribe,evalu,johfa,jonlu,mikki
```

```
TDBD30:*:9194:diasz,snt,ulfni
```

```
TDDA13:*:9161:bouka,patdo
```

```
HKGD20:*:9436:bjojo,erih,jonlu,matar,yuhhu
```

Group  
Name

Group ID

Members

Password

# A Look at Linux: Processes

- All execution takes place in processes
  - Each process may consist of several threads
  - Every process has its own (protected) address space
  - Every process has an ID, a parent, and a controlling tty
  - Processes have a state (running, stopped, suspended, etc)
  
- Processes can communicate
  - Signals are simple asynchronous messages
  - Processes can share memory areas
  - Processes can communicate using pipes
  - Processes can communicate using sockets

# Example of processes

```
% ps -H -eo s,pid,ppid,tty,user,cmd
S  PID  PPID TT      USER      CMD
S   1    0 ?      root      init [2]
S 2188 2194 ?      snmp      /usr/sbin/snmpd -Lsd -Lf /dev/null -u snmp
S 2194 2194 ?      root      /usr/sbin/sshd
S 24294 2194 ?      root      sshd: davby [priv]
S 24296 24294 ?      davby      sshd: davby@pts/0
S 24297 24296 pts/0      davby      -sh
R 24304 24297 pts/0      davby      ps -H -eo s,pid,ppid,tty,user,cmd
S 2206 1 ?      uml-net     /usr/bin/uml_switch -tap tap0 -unix
S 2273 1 ?      statd      /sbin/rpc.statd
S 2297 1 ?      root       sendmail: MTA: accepting connections
S 2323 1 ?      ntp        /usr/sbin/ntp -p /var/run/ntp.pid
S 2333 1 ?      daemon     /usr/sbin/atd
```

Process ID

Parent process ID

Controlling terminal

EUID

Command

# Signals

- User point-of-view: suspend, resume, kill processes

```
% ps axu | grep '[e]macs'  
davby 24613 0.6 0.2 9604 4596 pts/1 S+ 15:47 0:00 emacs -nw  
% kill -HUP 24613  
% ps axu | grep '[e]macs'  
%
```

- Send arbitrary signals using kill command
- If typing directly to process's controlling tty
  - C-c sends INTR
  - C-z sends TSTP
  - C-\ sends QUIT



# Privilege elevation

- Users gain extra privileges by changing EUID or starting processes with a different EUID than the current one

```
% ps -H -eo s,pid,ppid,tty,user,cmd
S   PID  PPID TT      USER      CMD
S     1    0 ?      root      init [2]
S  2194    1 ?      root      /usr/sbin/sshd
S  24294  2194 ?      root      sshd: davby [priv]
S  24296  24294 ?      davby    sshd: davby@pts/0
S  24297  24296 pts/0    davby    -sh
R  24321  24297 pts/0    davby
S  24312  2194 ?      root      sshd: davby [priv]
S  24314  24312 ?      davby    sshd: davby@pts/1
S  24315  24314 pts/1    davby    -sh
S  24319  24315 pts/1    root      passwd
```

sshd changed EUID from root to davby

passwd being run by davby with EUID root

# How does privilege elevation work?

- Programs can change their own EUID/EGID
  - The seteuid system call changes the EUID
  - The setegid system call changes the EGID
  - Very strict limitations on who can change to what
  
- Programs can have the setuid/setgid bits set
  - When setuid program started, process assumes file owner as EUID
  - When setgid program started, process assumes file group as EGID

# Example of setuid/setgid programs

```
% ls -l passwd crontab mail
-rwxr-sr-x 1 root news 26380 Dec 20 2006 crontab
-rwsr-xr-x 1 root root 28480 Feb 27 08:53 passwd
-rwsr-sr-x 1 root mail 72544 Apr 30 2006 procmail
%
```

crontab is setgid news

passwd is setuid root

procmail is setuid root and setgid mail

# The shell

- When a user logs in, the login program starts a shell
- The shell accepts and interprets commands from the user
  - Handles I/O redirection, environment variables, etc
- Two kinds of commands: built-in and external
  - Built-in: affect the shell itself (e.g. cd) or are run often (e.g. echo)
  - External: most everything else
  - Also: programming structures (e.g. if-then-else)
- External commands are just files with the execute permission set that are in a directory listed in the PATH variable

# System startup

- What happens when you start Linux
  1. The computer firmware (BIOS) loads the boot loader
  2. The boot loader loads and executes the operating system
  3. The operating system runs the /sbin/init program
  4. The /sbin/init program does what /etc/inittab says to do

```
% cat /etc/inittab  
id:2:initdefault:  
si::sysinit:/etc/init.d/rcS  
11:1:wait:/etc/init.d/rc 1  
12:2:wait:/etc/init.d/rc 2  
13:3:wait:/etc/init.d/rc 3  
1:23:respawn:/sbin/getty 38400 tty1  
2:23:respawn:/sbin/getty 38400 tty2  
3:23:respawn:/sbin/getty 38400 tty3  
%
```

Set default run level to 2

To do when initializing system

To do (once) when entering run level 2

To do when entering run level 2 or 3 (and when process terminates it is restarted)



# Typical (sysvinit) system startup

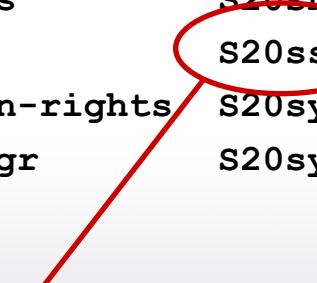
- /etc/init.d/rc script is run with run level as argument
  - Runs scripts in /etc/rcN.d directory
  - Scripts that start with K are run with argument **stop**
  - Scripts that start with S are run with argument **start**
- Note that /etc/rcS.d scripts are also run during boot

```
% ls /etc/rc2.d
```

S10sysklogd	S11klogd	S18portmap	S19autofs	%
-------------	----------	------------	-----------	---

S20cupsys	S20dbus	S20devtun-rights	S20dirmngr
S20snmpd	S20ssh	S20sysfsutils	S20sysinfo

S21exim	S23ntp	S89atd	S89cron
---------	--------	--------	---------



Run with argument **start** to start ssh service

# System logging

- Linux (and unix) systems and software are often very chatty
  - Detailed messages often show up in system and application logs
  - By default, log files are stored in /var/log

```
% ls -F /var/log
account/          cfengine.log.0      fsck           mail.log.0      quagga/
acpid            cfengine.log.1.gz   kern.log       mail.log.1.gz  samba/
acpid.1.gz       daemon.log        kern.log.0    mail.warn      syslog
apache/          daemon.log.0     kern.log.1.gz  mail.warn.0    syslog.0
apache2/         daemon.log.1.gz   ksymoops/     mail.warn.1.gz syslog.1.gz
aptitude         debug           lastlog        messages      user.log
aptitude.1.gz    debug.0         lpr.log        messages.0    user.log.0
auth.log         debug.1.gz      mail.err       messages.1.gz user.log.1.gz
auth.log.0       dmesg          mail.err.0    mysql/        uucp.log
auth.log.1.gz    dpkg.log        mail.err.1.gz  nagios/       vtund/
boot             dpkg.log.1     mail.info      nessus/       wtmp
btmp             fai/           mail.info.0   news/        wtmp.1
btmp.1          faillog        mail.info.1.gz ntpstats/    wtmp.report
cfengine.log     fontconfig.log  mail.log       pycentral.log %
%
```

# A Look at Linux: System logging

- Controlled by system logging service (syslogd)
  - Configured in /etc/syslog.conf
  - Controls what information goes where
  - Controls what **level** of information is logged

```
% head -2 /var/log/syslog.conf
```

```
auth,authpriv *
```

```
*.*;auth,authpriv.none
```

```
/var/log/auth.log
```

```
-/var/log/syslog
```

Log everything except auth messages to /var/log/syslog

# System logging

```
% /etc/init.d/bind9 reload
% tail -12 /var/log/syslog
Aug 20 08:29:22 sysinst-gw postfix/cleanup[26219]: 3FE26748B9: message-
id=<20070820062922.3FE26748B9@sysinst-gw.sysinst.ida.liu.se>
Aug 20 08:29:22 sysinst-gw postfix/bounce[26258]: F24931F4B7: sender non-
delivery notification: 3FE26748B9
Aug 20 08:29:22 sysinst-gw postfix/local[26259]: 3FE26748B9: to=<nagios@sysinst-
gw.sysinst.ida.liu.se>, relay=local, delay=0.16, delays=0.05/0.04/0/0.07,
dsn=2.0.0, status=sent (delivered to mailbox)
Aug 20 09:00:54 sysinst-gw named[7673]: loading configuration from
'/etc/bind/named.conf'
Aug 20 09:00:54 sysinst-gw named[7673]: zone 189.236.130.in-addr.arpa/IN: loaded
serial 2007081500
Aug 20 09:00:54 sysinst-gw named[7673]: zone 189.236.130.in-addr.arpa/IN:
sending notifies (serial 2007081500)
Aug 20 09:00:54 sysinst-gw named[7673]: zone sysinst.ida.liu.se/IN: loaded
serial 2007081500
Aug 20 09:00:54 sysinst-gw named[7673]: zone sysinst.ida.liu.se/IN: sending
notifies (serial 2007081500)
Aug 20 09:00:54 sysinst-gw named[7673]: client 130.236.177.25#34505: transfer of
'189.236.130.in-addr.arpa/IN': AXFR-style IXFR started
Aug 20 09:00:54 sysinst-gw named[7673]: client 130.236.177.25#34505: transfer of
'189.236.130.in-addr.arpa/IN': AXFR-style IXFR ended
```

# Where to next?

- Any basic textbook on Linux will teach you more
- The UML lab will familiarize you with the virtual environment you will be using (and in the virtual environment you root access, so you can do anything you want)
- The LXB lab will go through a large number of basic Linux issues you need when installing and managing systems

